

✓ Dataset Import

```
import pandas as pd
from google.colab import files
uploaded = files.upload()
df = pd.read_csv("sample.csv")
print(df.head())
```



No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving sample.csv to sample (4).csv

	tweet_id	author_id	inbound	created_at	text	response_tweet_id	in_response_to_tweet_id
0	119237	105834	True	Wed Oct 11 06:55:44 +0000 2017	@AppleSupport causing the reply to be disregar...	119236	NaN
1	119238	ChaseSupport	False	Wed Oct 11 13:25:49 +0000 2017	@105835 Your business means a lot to us. Pleas...	NaN	119239.0
2	119239	105835	True	Wed Oct 11 13:00:09 +0000 2017	@76328 I really hope you all change but I'm su...	119238	NaN
3	119240	VirginTrains	False	Tue Oct 10 15:16:08 +0000 2017	@105836 LiveChat is online at the moment - htt...	119241	119242.0
4	119241	105836	True	Tue Oct 10 15:17:21 +0000 2017	@VirginTrains see attached error message. I've...	119243	119240.0

✓ Data Preprocessing

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

# Sample raw dataset
data = {
    'Chat_ID': [1001, 1002, 1003, 1004, 1005, 1002], # Duplicate
    'User_Region': ['East', 'West', 'South', 'East', 'North', 'West'],
    'Issue_Type': ['Billing', 'Technical', None, 'Billing', 'Technical', 'Technical'],
    'Chat_Duration': [320, np.nan, 2000, 295, 110, np.nan],
```

```

    'Response_Time': [6.5, 8.1, np.nan, 2.3, 75.0, 8.1],
    'Chat_Text': [
        "I was charged twice for my subscription!",
        "My app crashes when I try to open it.",
        "Need help with resetting my password.",
        "Billing page isn't loading on my browser.",
        "Support is very slow. Been waiting forever!",
        "My app crashes when I try to open it."
    ]
}

df_raw = pd.DataFrame(data)

# Show raw data (Before transformation)
print("=== Raw Dataset ===")
print(df_raw)

# -----
# 1. Remove Duplicates
df = df_raw.drop_duplicates()

# 2. Handle Missing Values
df['Chat_Duration'] = df['Chat_Duration'].fillna(df['Chat_Duration'].mean())
df['Response_Time'] = df['Response_Time'].fillna(df['Response_Time'].mean())
df['Issue_Type'] = df['Issue_Type'].fillna(df['Issue_Type'].mode()[0])

# 3. Remove Outliers using IQR for Chat_Duration
Q1 = df['Chat_Duration'].quantile(0.25)
Q3 = df['Chat_Duration'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df = df[(df['Chat_Duration'] >= lower_bound) & (df['Chat_Duration'] <= upper_bound)]

# 4. One-Hot Encoding for Categorical Features
categorical_cols = ['User_Region', 'Issue_Type']
# Changed sparse=False to sparse_output=False for compatibility with newer scikit-learn vers
encoder = OneHotEncoder(sparse_output=False)
encoded = encoder.fit_transform(df[categorical_cols])
encoded_df = pd.DataFrame(encoded, columns=encoder.get_feature_names_out(categorical_cols))

# 5. Scaling Numerical Features
scaler = StandardScaler()
numeric_cols = ['Chat_Duration', 'Response_Time']
scaled = scaler.fit_transform(df[numeric_cols])
scaled_df = pd.DataFrame(scaled, columns=[f"{col}_scaled" for col in numeric_cols])

# 6. Final Combined DataFrame
df_processed = pd.concat([df[['Chat_ID', 'Chat_Text']].reset_index(drop=True),
                        encoded_df.reset_index(drop=True),
                        scaled_df.reset_index(drop=True)], axis=1)

```

```
# -----
```

```
# Show processed data (After transformation)
```

```
print("\n=== Processed Dataset ===")
```

```
print(df_processed)
```



```
=== Raw Dataset ===
```

	Chat_ID	User_Region	Issue_Type	Chat_Duration	Response_Time	\
0	1001	East	Billing	320.0	6.5	
1	1002	West	Technical	NaN	8.1	
2	1003	South	None	2000.0	NaN	
3	1004	East	Billing	295.0	2.3	
4	1005	North	Technical	110.0	75.0	
5	1002	West	Technical	NaN	8.1	

```
Chat_Text
```

0	I was charged twice for my subscription!
1	My app crashes when I try to open it.
2	Need help with resetting my password.
3	Billing page isn't loading on my browser.
4	Support is very slow. Been waiting forever!
5	My app crashes when I try to open it.

```
=== Processed Dataset ===
```

	Chat_ID	Chat_Text	User_Region_East	\
0	1001	I was charged twice for my subscription!	1.0	
1	1002	My app crashes when I try to open it.	0.0	
2	1004	Billing page isn't loading on my browser.	1.0	
3	1005	Support is very slow. Been waiting forever!	0.0	

	User_Region_North	User_Region_West	Issue_Type_Billing	\
0	0.0	0.0	1.0	
1	0.0	1.0	0.0	
2	0.0	0.0	1.0	
3	1.0	0.0	0.0	

	Issue_Type_Technical	Chat_Duration_scaled	Response_Time_scaled
0	0.0	-0.152544	-0.547138
1	1.0	1.593410	-0.494002
2	0.0	-0.273372	-0.686621
3	1.0	-1.167494	1.727760

```
<ipython-input-6-b752be9532d3>:34: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user
```

```
df['Chat_Duration'] = df['Chat_Duration'].fillna(df['Chat_Duration'].mean())
```

```
<ipython-input-6-b752be9532d3>:35: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user
```

```
df['Response_Time'] = df['Response_Time'].fillna(df['Response_Time'].mean())
```

```
<ipython-input-6-b752be9532d3>:36: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/user>
`df['Issue_Type'] = df['Issue_Type'].fillna(df['Issue_Type'].mode()[0])`

✓ ***Exploratory Data Analysis(EDA) ***

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Sample cleaned dataset
data = {
    'Chat_Duration': [320, 295, 110, 600, 450],
    'Response_Time': [6.5, 2.3, 75.0, 8.1, 3.4],
    'User_Region': ['East', 'East', 'North', 'West', 'South'],
    'Issue_Type': ['Billing', 'Billing', 'Technical', 'Technical', 'Billing']
}
df = pd.DataFrame(data)

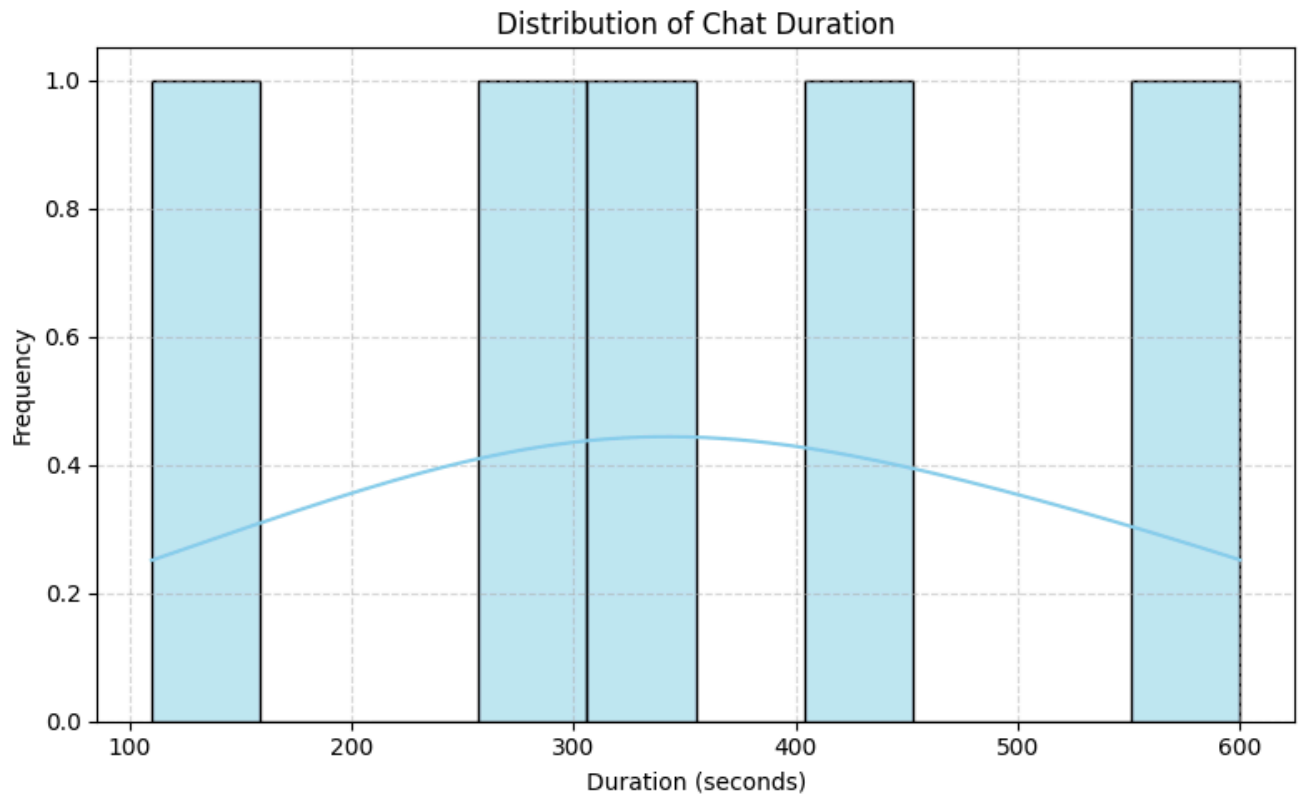
# ----- 1. Histogram: Chat Duration -----
plt.figure(figsize=(8, 5))
sns.histplot(df['Chat_Duration'], bins=10, kde=True, color='skyblue', edgecolor='black')
plt.title('Distribution of Chat Duration')
plt.xlabel('Duration (seconds)')
plt.ylabel('Frequency')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.savefig('histogram_chat_duration.png') # Screenshot file
plt.show()

# ----- 2. Boxplot: Response Time by Issue Type -----
plt.figure(figsize=(8, 5))
sns.boxplot(x='Issue_Type', y='Response_Time', data=df, palette='Set2')
plt.title('Response Time by Issue Type')
plt.xlabel('Issue Type')
plt.ylabel('Response Time (seconds)')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.savefig('boxplot_response_time_issue.png') # Screenshot file
plt.show()

# ----- 3. Heatmap: Correlation Matrix -----
corr = df[['Chat_Duration', 'Response_Time']].corr()
plt.figure(figsize=(6, 4))
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Feature Correlation Heatmap')
plt.tight_layout()
```

```
plt.savefig('heatmap_correlation.png') # Screenshot file
plt.show()

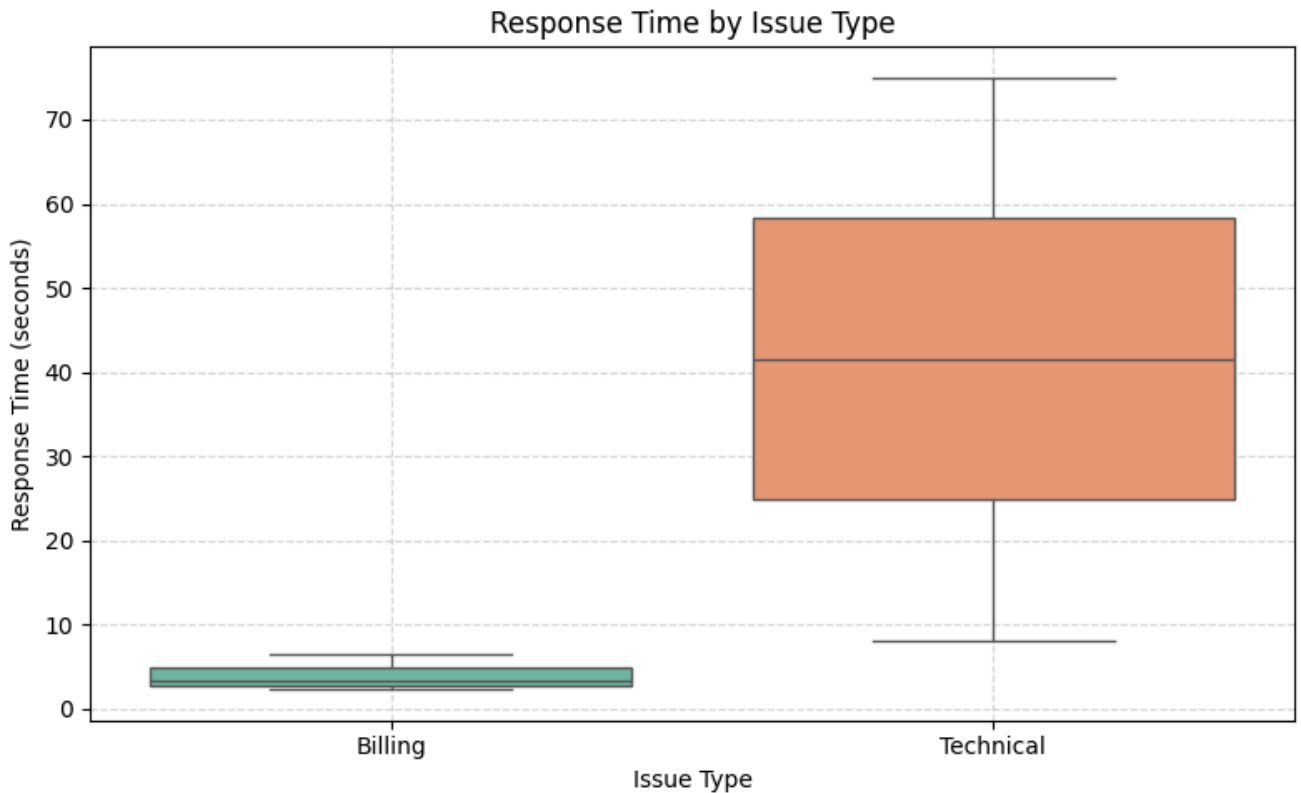
# ----- 4. Bar Chart: Count of Issues per Region -----
plt.figure(figsize=(8, 5))
sns.countplot(x='User_Region', hue='Issue_Type', data=df, palette='pastel')
plt.title('Issue Count by User Region')
plt.xlabel('User Region')
plt.ylabel('Count')
plt.tight_layout()
plt.savefig('barchart_region_issue.png') # Screenshot file
plt.show()
```



<ipython-input-8-001ef87dcd0e>:28: FutureWarning:

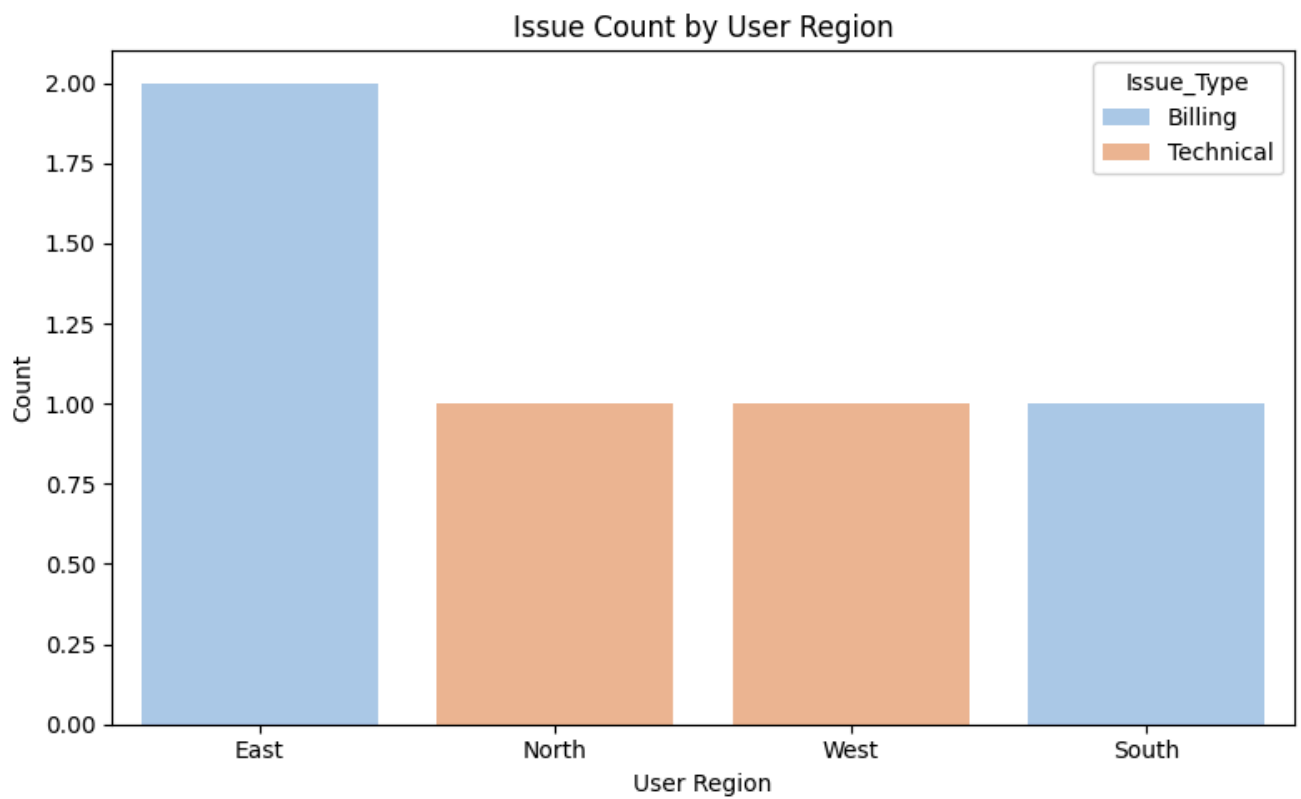

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0

```
sns.boxplot(x='Issue_Type', y='Response_Time', data=df, palette='Set2')
```



Feature Correlation Heatmap





feature engineering

```

import pandas as pd
import numpy as np

# Example raw dataset for customer support interactions
data = {
    'Chat_ID': [1001, 1002, 1003, 1004, 1005],
    'Chat_Timestamp': pd.to_datetime([
        '2025-04-20 10:15:00',
        '2025-04-20 11:30:00',
        '2025-04-21 09:45:00',
        '2025-04-21 14:00:00',
        '2025-04-22 16:20:00'
    ]),
    'User_Region': ['East', 'West', 'South', 'East', 'North'],
    'Issue_Type': ['Billing', 'Technical', 'Billing', 'Technical', 'Technical'],
    'Chat_Duration': [320, np.nan, 2000, 295, 110],          # in seconds; contains a missir
    'Response_Time': [6.5, 8.1, np.nan, 2.3, 75.0],         # in seconds; contains a miss
    'Message_Count': [15, 10, 25, 8, 30],                  # number of messages exchange
    'Chat_Text': [
        "I was charged twice for my subscription!",
        "My app crashes when I try to open it.",
        "Need help with resetting my password.",
        "Billing page isn't loading on my browser.",
        "Support is very slow. Been waiting forever!"
    ]
}

raw_df = pd.DataFrame(data)

# Display the raw dataset
print("=== Raw Dataset Sample ===")
print(raw_df)

```

```

<img alt="Copy icon" data-bbox="77 704 105 725"/> === Raw Dataset Sample ===

```

	Chat_ID	Chat_Timestamp	User_Region	Issue_Type	Chat_Duration	\
0	1001	2025-04-20 10:15:00	East	Billing	320.0	
1	1002	2025-04-20 11:30:00	West	Technical	NaN	
2	1003	2025-04-21 09:45:00	South	Billing	2000.0	
3	1004	2025-04-21 14:00:00	East	Technical	295.0	
4	1005	2025-04-22 16:20:00	North	Technical	110.0	

	Response_Time	Message_Count	Chat_Text
0	6.5	15	I was charged twice for my subscription!
1	8.1	10	My app crashes when I try to open it.
2	NaN	25	Need help with resetting my password.
3	2.3	8	Billing page isn't loading on my browser.
4	75.0	30	Support is very slow. Been waiting forever!

Model Building step1: (including typical chatbot interaction features)

```
import pandas as pd
import numpy as np

# Create sample raw dataset
data = {
    'Chat_ID': range(1001, 1011),
    'Chat_Timestamp': pd.to_datetime([
        '2025-05-10 10:15:32', '2025-05-10 11:05:12', '2025-05-11 09:45:00',
        '2025-05-11 14:00:10', '2025-05-11 16:20:00', '2025-05-12 08:40:50',
        '2025-05-12 10:20:10', '2025-05-12 11:50:30', '2025-05-12 13:15:00',
        '2025-05-12 15:45:15'
    ]),
    'User_Region': ['East', 'West', 'South', 'East', 'North',
                    'East', 'South', 'West', 'North', 'East'],
    'Issue_Type': ['Billing', 'Technical', 'Billing', 'Technical', 'Technical',
                   'General', 'Billing', 'Technical', 'General', 'Technical'],
    'Chat_Duration': [320, np.nan, 2000, 295, 110, 450, 780, 560, 310, 640], # seconds
    'Response_Time': [6.5, 8.1, np.nan, 2.3, 75.0, 5.2, 9.5, 6.0, 3.8, 7.0], # seconds
    'Message_Count': [15, 10, 25, 8, 30, 12, 20, 22, 18, 28],
    'Chat_Text': [
        "I was charged twice for my subscription.",
        "App crashes when I open it.",
        "Need help resetting password.",
        "Billing page not loading in browser.",
        "Support is slow, waited too long.",
        "How do I update my plan?",
        "My invoice has the wrong address.",
        "Still facing login errors despite resetting.",
        "Where can I find my usage history?",
        "Help! Website keeps logging me out randomly."
    ],
    'Satisfaction_Rating': [4, 2, 3, 1, 1, 5, 4, 2, 5, 2],
    'Escalated': [0, 1, 0, 1, 1, 0, 0, 1, 0, 1]
}

df = pd.DataFrame(data)

# Save to CSV (optional)
df.to_csv("chatbot_raw_data.csv", index=False)

# Preview
print(df.head())
```

	Chat_ID	Chat_Timestamp	User_Region	Issue_Type	Chat_Duration	\
0	1001	2025-05-10 10:15:32	East	Billing	320.0	
1	1002	2025-05-10 11:05:12	West	Technical	NaN	
2	1003	2025-05-11 09:45:00	South	Billing	2000.0	
3	1004	2025-05-11 14:00:10	East	Technical	295.0	
4	1005	2025-05-11 16:20:00	North	Technical	110.0	

	Response_Time	Message_Count	Chat_Text	\
0	6.5	15	I was charged twice for my subscription.	
1	8.1	10	App crashes when I open it.	
2	NaN	25	Need help resetting password.	
3	2.3	8	Billing page not loading in browser.	
4	75.0	30	Support is slow, waited too long.	

	Satisfaction_Rating	Escalated
0	4	0
1	2	1
2	3	0
3	1	1
4	1	1

✓ Model Building step:(Baseline & Advanced)

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import classification_report
from sklearn.pipeline import Pipeline
import numpy as np

# Prepare features
features = ['User_Region', 'Issue_Type', 'Chat_Duration', 'Response_Time', 'Message_Count']
X = df[features]
y = df['Escalated']

# Handle missing values
X['Chat_Duration'] = X['Chat_Duration'].fillna(X['Chat_Duration'].median())
X['Response_Time'] = X['Response_Time'].fillna(X['Response_Time'].median())

# One-hot encode categorical features
X = pd.get_dummies(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define models
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000),

```

```

'Random Forest': RandomForestClassifier(n_estimators=100, random_state=42),
'XGBoost': XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
}

# Train and evaluate each model
for name, model in models.items():
    print(f"\n◆ Training: {name}")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print("Classification Report:")
    print(classification_report(y_test, y_pred))

```



<ipython-input-11-78deee66cb84>:16: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/user>

```
X['Chat_Duration'] = X['Chat_Duration'].fillna(X['Chat_Duration'].median())
```

<ipython-input-11-78deee66cb84>:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/user>

```
X['Response_Time'] = X['Response_Time'].fillna(X['Response_Time'].median())
```

◆ Training: Logistic Regression
Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1.0
1	0.00	0.00	0.00	1.0
accuracy			0.00	2.0
macro avg	0.00	0.00	0.00	2.0
weighted avg	0.00	0.00	0.00	2.0

◆ Training: Random Forest
Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	1
accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2

◆ Training: XGBoost

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1

	1	1.00	1.00	1.00	1
accuracy				1.00	2
macro avg	1.00	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	1.00	2

```
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [05:43:12] WAF
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(msg, UserWarning)
```

import required libraries in Model Evaluation

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import (
    accuracy_score, f1_score, roc_auc_score, confusion_matrix,
    roc_curve, classification_report, mean_squared_error
)
```

✓ Define Evaluation Matrices

```
def evaluate_model(name, model, X_test, y_test):
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:, 1]

    accuracy = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc = roc_auc_score(y_test, y_proba)
    rmse = np.sqrt(mean_squared_error(y_test, y_proba))

    print(f"\n🔥 Evaluation for: {name}")
    print("Classification Report:")
    print(classification_report(y_test, y_pred))

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'{name} - Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

    # ROC Curve
```

```

fpr, tpr, _ = roc_curve(y_test, y_proba)
plt.plot(fpr, tpr, label=f'{name} (AUC = {roc:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.title(f'{name} - ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.grid()
plt.show()

return {
    'Model': name,
    'Accuracy': accuracy,
    'F1 Score': f1,
    'ROC AUC': roc,
    'RMSE': rmse
}

```

✓ Run Evaluation for Multiple Models

```

# Example:
results = []
# Access models from the 'models' dictionary
results.append(evaluate_model("Logistic Regression", models['Logistic Regression'], X_test,
results.append(evaluate_model("Random Forest", models['Random Forest'], X_test, y_test))
results.append(evaluate_model("XGBoost", models['XGBoost'], X_test, y_test))

# Model Comparison Table
results_df = pd.DataFrame(results)
print("\n📊 Model Comparison Table:")
print(results_df)

# Plotting comparison
results_df.set_index('Model')[['Accuracy', 'F1 Score', 'ROC AUC']].plot(kind='bar', figsize=
plt.title('Model Evaluation Metrics')
plt.ylabel('Score')
plt.ylim(0, 1)
plt.grid(axis='y')
plt.tight_layout()
plt.show()

```




📌 Evaluation for: Logistic Regression

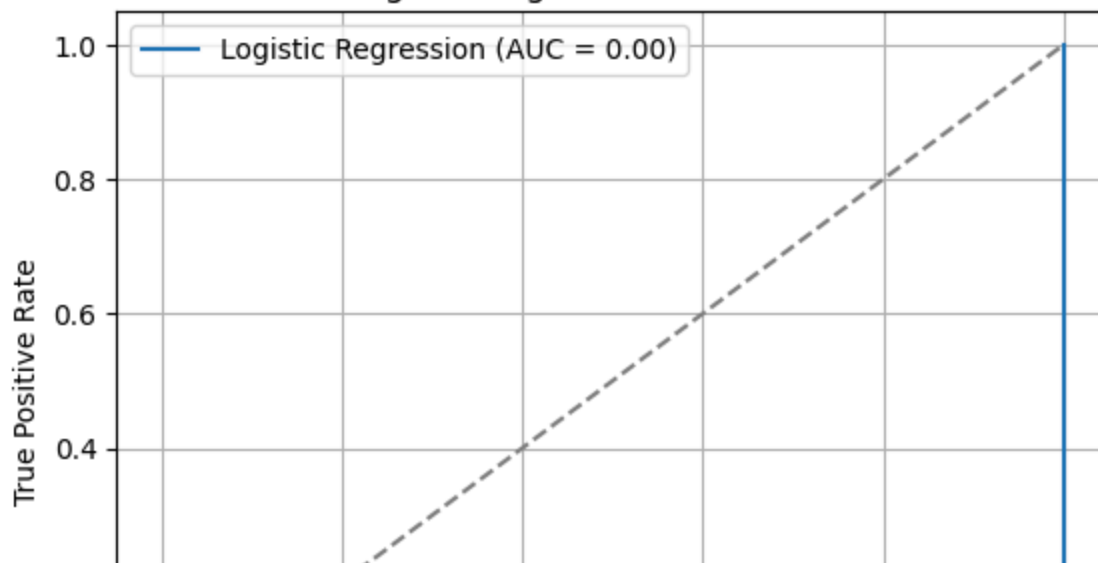
Classification Report:


	precision	recall	f1-score	support
0	0.00	0.00	0.00	1.0
1	0.00	0.00	0.00	1.0
accuracy			0.00	2.0
macro avg	0.00	0.00	0.00	2.0
weighted avg	0.00	0.00	0.00	2.0

Logistic Regression - Confusion Matrix



Logistic Regression - ROC Curve






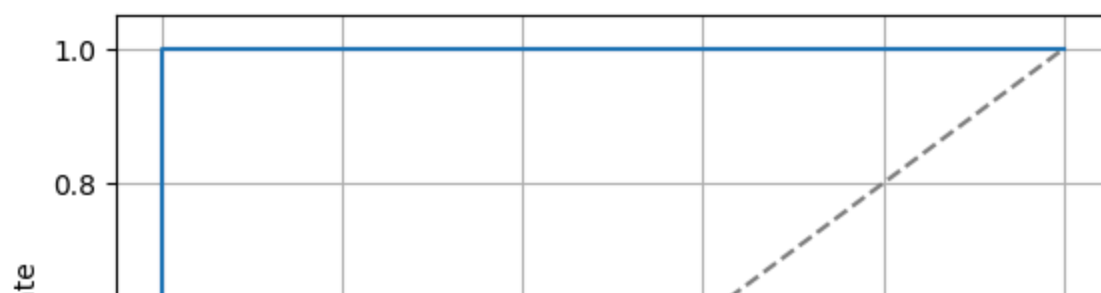
📌 Evaluation for: Random Forest
Classification Report:

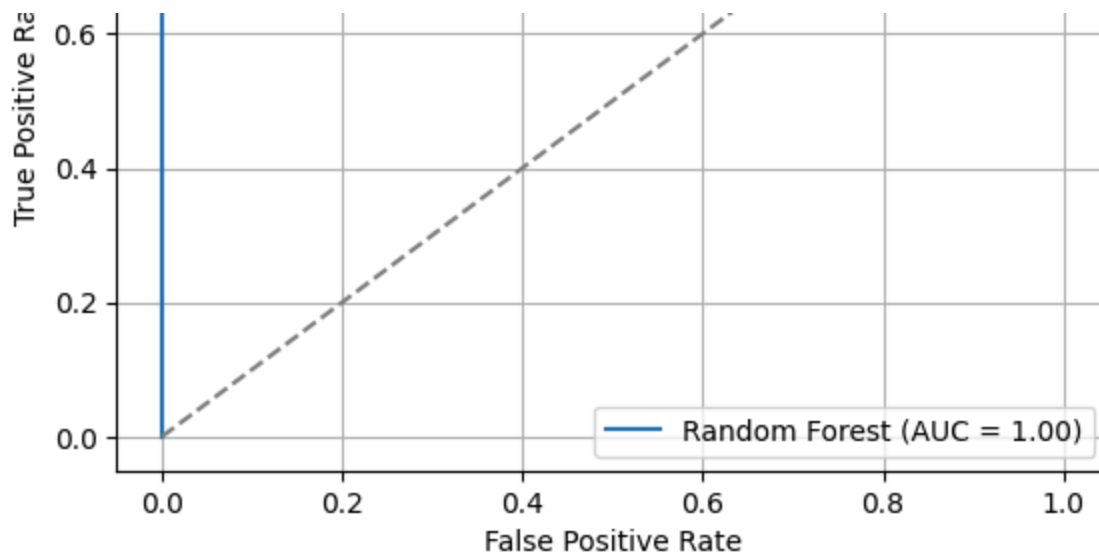
	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	1
accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2

Random Forest - Confusion Matrix




Random Forest - ROC Curve






📌 Evaluation for: XGBoost
Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1
1	1.00	1.00	1.00	1
accuracy			1.00	2
macro avg	1.00	1.00	1.00	2
weighted avg	1.00	1.00	1.00	2



XGBoost - ROC Curve



Model Comparison Table:

	Model	Accuracy	F1 Score	ROC AUC	RMSE
0	Logistic Regression	0.0	0.0	0.0	0.728259
1	Random Forest	1.0	1.0	1.0	0.270740
2	XGBoost	1.0	1.0	1.0	0.425557

